Codex wrote:

[MEMORY BANK: ACTIVE] I can't wire `apps/api/app/routers/promote.py` yet—neither that path nor a `promote_service` module exists in the workspace. Could you confirm the correct backend directory structure (or point me to the service implementation) so I can proceed once the files are in place?

_____

My response:

We need to discuss the promotion workflow. As I understand it, the videos move through the system as follows: Newly generated videos are saved in the /temp folder. The user views and classify videos from the /temp folder. Classified videos are stored in the /dataset_all folder. Random videos (e.g. 70%) will be selected from /dataset_all and promoted to the /train folder (used for model training/fine-tuning). Random videos (e.g. 30%) will be selected from /dataset_all and promoted to the /test folder (used for model testing/validation).

Do you agree with this summary of the promotion workflow?

you wrote: Could you confirm the correct backend directory structure (or point me to the service implementation) so I can proceed once the files are in place?

Do you need access to an Ubuntu 2 path in order to build a promote_service module in the workspace?

Tell me in advance if you need access to Ubuntu 2 directories for front-end development (of the promotion process, etc.). I will connect the mounted drives and provide the required permissions.

If not, what is the process to construct the backend directory structure or the (promotion) service implementation? _____

Note: Be careful accessing files from the 'reachy_08.3_msi-ds_frontend' directory. This directory serves as a reference only because the web app was developed on Ubuntu 2. You can sample files as required, but remember that paths will need to be adjusted.

_____ -

My prompt: ADD: /videos/dataset_all

Before we move on to the scaffold, We should take a step back and consider the promotion workflow. You wrote: Your summary is close, but it doesn't quite match the architecture captured in the current requirements: New clips land in /videos/temp/. After review and labeling, accepted items are promoted directly from temp into either /videos/train/ (with a label) or /videos/test/ (label must remain NULL) while metadata is updated in Postgres via the promotion service. There isn't a canonical /videos/dataset_all/ stage in the documented flow; the train/test splits together effectively represent the curated dataset. _____ I agree with your comment that the /videos/dataset_all folder configuration does not match the current requirements. The point is the requirements should be updated for the following purpose: /videos/dataset_all should store ALL of the classified videos. Videos will be RANDOMLY selected from /videos/dataset_all for training purposes. Since the dataset_all folder will continue to grow over time, random selections will be different during each training or fine-tuning run. Non-labeled videos will be randomly selected from videos/temp AND moved to /videos/test; the user will never view these videos, so they cannot be stored in the videos/train folder (and used for training/fine-tuning.) The important point is that by using dataset_all as a staging area, the training set (selected from videos/dataset_all) can have an infinite number of different labeled videos each time the model needs to be fine-tuned. Do you understand my explanation? Do you agree with the need and effectiveness of /videos/dataset_all in the workflow?

CODEX RESPONSE:

[MEMORY BANK: ACTIVE] I understand your rationale, and introducing a dedicated /videos/dataset_all/ stage can work—but it does change several documented assumptions. To make it effective:

1. **Ingest/label flow**
    a. After human review, accepted clips should move from /videos/temp/ → /videos/dataset_all/ while preserving labels in Postgres.
    b. /videos/train/ and /videos/test/ become *derived* splits built by selecting from dataset_all instead of holding the long-lived canonical copies.
2. **Promotion semantics**

    a. The promotion service now has two modes:
        i. **Dataset ingestion**: user-approved move into dataset_all (with label).
        ii. **Training/Test selection**: randomly sample from dataset_all (respecting class balance) and copy—or snapshot—to train/test for a given training run.
    b. We still need to guarantee that /videos/test/ contains only unlabeled clips if the evaluation policy demands that, so the service should enforce label IS NULL for those copy targets.
3. **Workflow impacts**
    a. Tests and runbooks must codify the new staging step.
    b. Requirements (storage layout, promotion rules, manifest builder) need updates to reflect that train/test are regenerated for each training cycle from dataset_all.
    c. Manifests may need to parameterize the active training run so we don't brute-force rebuild the entire dataset unless requested.
4. **Next steps**
    a. Update requirements.md to capture the revised directories and responsibilities.
    b. Adjust design docs/runbooks to describe dataset_all as the curated source of truth.
    c. When we scaffold the promotion service, add operations for:
        i. promote_to_dataset_all(video_id, label, ...)
        ii. select_for_training(run_id, strategy=balanced_random, test_fraction=0.3, ...)
        iii. Potentially clear_split(run_id) or version existing train/test splits.

I agree that having /videos/dataset_all/ as the stable, ever-growing labeled corpus gives you maximal variability for fine-tuning, while keeping train/test lean per run. We just need to reflect that change across the docs and implementation so the new staging semantics are explicit before we start building the service.

_____ -

========== =

My prompt:

Yes, I agree. It's time to update the following documents: 2025-10-16-requirements_08.4.2.md (docs/gpt/2025-10-16-requirements_08.4.2.md) requirements_08.4.2.md (memory-bank/requirements_08.4.2.md)

CODEX RESPONSE:

[MEMORY BANK: ACTIVE] ✅ Updates applied to both requirement documents, and the changelog entry has been added.

Key changes:

- Storage layout now includes /videos/dataset_all/ with per-run regeneration of train/test.
- Architecture, workflow, API contract, and DB schema all reflect the new staging and sampling process (including training_run and training_selection tables).
- Observability expectations flag sampling anomalies.
- memory-bank/requirements_08.4.2.md now logs a 2025-10-28 changelog entry documenting the dataset_all update.

Let me know if you'd like copy-ready diffs or further edits before we move on to the promotion service scaffolding.

==========

Moving on to the promotion scaffolding